

Here's the overall assignment so you understand our objectives: "# Instructor-led Lab: Descriptive Statistics

In this assignment you will practice implementing statistical approaches in Python. You will use the datasets you previously made use of in the [Week 7 Independent lab](https://github.com/UM-BGEN632/week7labs/blob/main/07_Independent_Lab.md).

Context

Your administrator was pleased with your work in creating subsets of the hospital data. Recall, you belong to a team assigned to assess the condition of the healthcare system in California. You currently work in the Information Systems department for a consulting firm working with the state government agency that oversees the healthcare system in California. The data you used included two files (see the two tables for metadata):

- * Data for 61 hospitals in the file

[CaliforniaHospitalData.csv]

(/data/CaliforniaHospitalData.csv)

- * Personnel data containing employee information within the file [CaliforniaHospitalData_Personnel.txt]

(/data/CaliforniaHospitalData_Personnel.txt)

This first table provides the variables in the hospital data.

Variable	Description
----------	-------------

HospitalID	The primary key of each hospital
------------	----------------------------------

Name	The legal name of the hospital
------	--------------------------------

Zip	Zipcode where the hospital is located
-----	---------------------------------------

Website	The url for the hospital's website
---------	------------------------------------

TypeControl	Indicates the primary managing entity of the hospital
-------------	---

Teaching	Indicates teaching status
----------	---------------------------

DonorType	This field indicates the most prominent group of donors
-----------	---

NoFTE	Number of full-time employees registered at the hospital
-------	--

NetPatRev	Net patient revenue
-----------	---------------------

InOperExp	Estimate of the inpatient operating costs
-----------	---

OutOperExp	Estimate of the outpatient operating costs
------------	--

OperRev	Operating revenue of the hospital
---------	-----------------------------------

--	--

| OperInc | Operating Income is the operating revenue less the operating expenses |
| AvlBeds | The number of available beds in the hospital |

This second table provides the data for the personnel data.

Variable	Description
HospitalID	The foreign key of the hospital where position is held
Work_ID	Primary key of the personnel
LastName	The last name of the personnel
FirstName	First name of the personnel
Gender	Gender of the individual
PositionID	The foreign key for the position held
PositionTitle	The title of this position
Compensation	The annual amount the position is compensated for service
MaxTerm	The maximum number of years an individual can serve in this position
StartDate	The beginning of service for this position

Prep Data and Add A New Record

Your purpose for this assignment is to conduct a descriptive analysis on the data as a precursor to model building. You will use Python to perform your analysis. Merge the two data files. Remove the following columns of data:

- * Work_ID
- * PositionID
- * Website

Select one of the existing hospitals in the data and create a new position for yourself. Put in your first name and last name. Put today's date as the start date. Select one of the positions as shown in the table below and fill out the data accordingly. Fill in the rest of the columns as you choose. You should have one new row of data.

Output the DataFrame in your notebook.

PositionTitle	Compensation	MaxTerm
Regional Representative	46978	4
State Board Representative	89473	3
Acting Director	248904	8
Safety Inspection Member	23987	2

After merging the data files together, convert any date-

time columns into a datetime data type. Ensure object data (e.g., `TypeControl`, `Teaching`, `DonorType`, `Gender`) has the categorical data type.

Descriptive Statistics

Please produce the following summary statistics:

- * Provide a summary of the mean, median, minimum value, and maximum value for each numeric variable.
- * Provide summary statistics of your text variables.
- * Provide summary statistics of your categorical variables.

Create some visuals to explore the data:

- * Create histograms of the following variables: `NoFTE`, `NetPatRev`, `InOperExp`, `OutOperExp`, `OperRev`, `OperInc`, and `AvlBeds`.
- * Create scatterplots of the following variables using `*net patient revenue*` as the target variable (i.e., place it on the `*y*-axis`): `NoFTE`, `InOperExp`, `OutOperExp`, `OperRev`, `OperInc`, and `AvlBeds`.
 - * Describe the linearity and trend of each plot.
- * Create a boxplot and assess the lack or presence of outliers for the following variables: `NoFTE`, `NetPatRev`, `InOperExp`, `OutOperExp`, `OperRev`, `OperInc`, and `AvlBeds`. You may need to split the variables up due to the range of values in the data (like we did in the tutorial).
- * Choose three variables and create a QQ plot, providing your assessment of the normality.
- * Perform a Shapiro-Wilk test on those three variables; do your test results coincide with the QQ plots?

****Save your notebook with output showing your code has been successfully executed and then submit for grading.****

Great—thanks for sharing the full scope. Here’s how I see our high-level roadmap:

1. Data preparation

- Read in both files
- Merge on `HospitalID`
- Drop `Work_ID`, `PositionID`, and `Website`
- Add your new personnel record with today’s date
- Convert date fields to datetime and object fields to categorical

2. Descriptive statistics

- Numeric summaries (mean, median, min, max)
- Text summaries
- Categorical summaries

3. Exploratory visuals

- Histograms for each numeric variable
- Scatterplots (each predictor vs. NetPatRev) with trend/linearity notes
- Boxplots to spot outliers
- QQ plots + Shapiro–Wilk tests for three chosen variables

4. Wrap-up

- Save notebook with outputs

Let me know which step you'd like to tackle first—e.g. setting up the data merge, adding your record, or jumping straight into summary stats—and I'll guide you through the code (or assemble cells) accordingly.



ok i need some help with my cell. My instructions are to
"## Prep Data and Add A New Record

Your purpose for this assignment is to conduct a descriptive analysis on the data as a precursor to model building. You will use Python to perform your analysis. Merge the two data files. Remove the following columns of data:

* Work_ID

* PositionID

* Website" My code is: "# load data and do a quick inspect (I copy-pasted this from my week 7 assignment)

```
CA_Hospital =
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData.csv")
```

```
CA_Hospital_Personnel =
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData_Personnel.txt", delimiter='\t')
merged_df = pd.merge(CA_Hospital,
CA_Hospital_Personnel, on='HospitalID')
```

```
print(CA_Hospital.columns)". That doesn't include all
```

the columns I should see, nor does it remove the three specified columns. Give me the correct cell, and explain my mistakes

Here's a single cell that will:

1. Read in both files
2. Merge on HospitalID
3. Drop Work_ID , PositionID , and Website
4. Show you the resulting columns and a quick preview

python

 Copy  Edit

```
import pandas as pd # 1. Load the data hospital_df = pd.read_csv(
"/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/" "My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData.csv" )
personnel_df = pd.read_csv(
"/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/" "My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData_Personnel.txt",
sep="\t" ) # 2. Merge on HospitalID merged_df = pd.merge(hospital_df,
personnel_df, on="HospitalID") # 3. Drop the unwanted columns merged_df
= merged_df.drop(columns=["Work_ID", "PositionID", "Website"]) # 4.
Inspect the result print("Columns after merge and drop:")
print(merged_df.columns.tolist()) merged_df.head()
```

What was going wrong in your original snippet:

- You printed `CA_Hospital.columns` , which only shows the hospital table's columns —so you never saw what the merged DataFrame contained.
- You never actually dropped the `Work_ID` , `PositionID` , or `Website` columns.
- The merge itself was fine, but without inspecting `merged_df.columns` you wouldn't know what to drop or whether the merge succeeded.
- By reassigning (`merged_df = merged_df.drop(...)`) you ensure those columns are removed before further work.

Let me know when you'd like to run this or move on to adding your new record!



next task is: "Select one of the existing hospitals in the data and create a new position for yourself. Put in your first name and last name. Put today's date as the start date. Select one of the positions as shown in the table below and fill out the data accordingly. Fill in the rest of the columns as you choose. You should have one new row of data.

Output the DataFrame in your notebook."

To add your own row you'll want to:

1. **Pick an existing HospitalID** (e.g. `merged_df["HospitalID"].iloc[0]` or any ID you like).
2. **Build a dict** containing every column in `merged_df` (except the ones you dropped), filling in your values:
 - `HospitalID` : the one you picked
 - `LastName / FirstName` : your name
 - `Gender` : whatever you choose
 - `PositionTitle` , `Compensation` , `MaxTerm` : from the table of four options
 - `StartDate` : today's date
 - For the remaining hospital-level columns (`Name` , `Zip` , `TypeControl` , etc.) you can either reuse the values from that hospital or invent reasonable stand-ins.
3. **Convert your dict to a one-row DataFrame** and append (or `pd.concat`) to `merged_df` .
4. **Display** `merged_df.tail()` (or `merged_df` fully) to confirm.

Here's a template of the code bits you need (not wrapped as a full cell):

python

 Copy  Edit

```
import pandas as pd # 1. choose an existing hospital existing_id =
merged_df["HospitalID"].iloc[0] # 2. grab its row as a dict to reuse
most hospital fields base = merged_df[merged_df["HospitalID"] ==
existing_id].iloc[0].to_dict() # 3. overwrite with your personnel info
base.update({ "LastName": "McKenzie", "FirstName": "Alex", "Gender":
"M", # or your choice "PositionTitle": "Regional Representative",
"Compensation": 46978, "MaxTerm": 4, "StartDate": pd.to_datetime("2025-
04-28") }) # 4. turn into 1-row DataFrame and append new_row_df =
pd.DataFrame([base]) merged_df = pd.concat([merged_df, new_row_df],
ignore_index=True) # 5. inspect merged_df.tail()
```

Key notes on mistakes to avoid

- You must include every column in your dict, not just the new ones—otherwise pandas will fill missing columns with NaN.
- Convert your "StartDate" to a datetime before appending, so later dtype conversions are smooth.
- Use `pd.concat(..., ignore_index=True)` rather than `.append()` , since `.append()` is deprecated.

Give that a shot and let me know if you hit any snags or want it wrapped into one cell!



```
next it says "| PositionTitle | Compensation | MaxTerm |
|:---|---:|:---:|
| Regional Representative | 46978 | 4 |
| State Board Representative | 89473 | 3 |
| Acting Director | 248904 | 8 |
| Safety Inspection Member | 23987 | 2 |
```

After merging the data files together, convert any date-time columns into a datetime data type. Ensure object data (e.g., `TypeControl`, `Teaching`, `DonorType`, `Gender`) has the categorical data type. "

Here's how you can convert your date-time column(s) and turn the specified object columns into categoricals without wrapping it into a full cell:

python

 Copy  Edit

```
# 1. Convert StartDate (and any other date columns) to datetime
merged_df['StartDate'] = pd.to_datetime(merged_df['StartDate']) # 2.
Convert object columns to categorical for col in ['TypeControl',
'Teaching', 'DonorType', 'Gender']: merged_df[col] =
merged_df[col].astype('category') # 3. Quick check of dtypes
merged_df.dtypes[['StartDate', 'TypeControl', 'Teaching', 'DonorType',
'Gender']]
```

What this does:

- `pd.to_datetime` safely parses your "StartDate" strings into `datetime64[ns]` .
- The loop over your list of column names casts each to `category` , which is more memory-efficient and better for downstream analysis.
- Finally, printing out those dtypes confirms the conversions.

why wouldn't Name, Last Name, First Name, Position Title, and Email columns also get converted from object to categorical?

Good question. It really comes down to the *semantics* and *cardinality* of those fields:

- **Name**, **LastName**, **FirstName**, and **Email** are effectively unique identifiers or near-unique strings. Turning them into pandas `category` would create as many categories as there are rows, offering no memory or performance benefit—and in fact can *slow* lookups when there's little repetition.
- **PositionTitle** does have repetition, but we didn't convert it in our minimal example simply because the assignment only called out certain "object" columns. If you

know you're going to group or filter by `PositionTitle` repeatedly, by all means cast it to categorical:

python



Copy



Edit

```
merged_df['PositionTitle'] =  
merged_df['PositionTitle'].astype('category')
```

In short, you convert columns to `category` when they represent a small, fixed set of possible values that you'll use for grouping, filtering, or modelling—otherwise it's fine (and often preferable) to leave free-text identifiers as plain `object` strings.



Next we have to "## Descriptive Statistics

Please produce the following summary statistics:

- * Provide a summary of the mean, median, minimum value, and maximum value for each numeric variable.
- * Provide summary statistics of your text variables.
- * Provide summary statistics of your categorical variables."

Here's how you can get each of those summaries (you can copy-paste these lines into your notebook):

python



Copy



Edit

```
import numpy as np # 1. Numeric variables: mean, median, min, max
num_cols = merged_df.select_dtypes(include=[np.number]).columns
numeric_summary = merged_df[num_cols].agg(['mean', 'median', 'min',
'max']).T print(numeric_summary) # 2. Text variables: count, unique,
top, freq # (these are the object-dtype columns that aren't coded as
categorical) text_cols = merged_df.select_dtypes(include=
['object']).columns.tolist() text_summary =
merged_df[text_cols].describe() print(text_summary) # 3. Categorical
variables: count, unique, top, freq cat_cols =
merged_df.select_dtypes(include=['category']).columns.tolist()
cat_summary = merged_df[cat_cols].describe() print(cat_summary)
```

What this does:

- **Numeric:** automatically picks up all number columns and computes exactly the four statistics you need.
- **Text:** `.describe()` on object dtypes gives you count, number of unique values, most frequent value ("top") and its frequency.
- **Categorical:** same `.describe()` but now on your explicitly cast categories.

Let me know if you'd like that wrapped into one full cell or any tweaks (e.g. excluding some columns from the "text" list).



Create some visuals to explore the data:

* Create histograms of the following variables: NoFTE, NetPatRev, InOperExp, OutOperExp, OperRev, OperInc, and AvlBeds.

It looks like `merged_df` isn't defined in the current cell environment. Let me prepend the cell with loading or referencing it—you can replace that snippet with your actual DataFrame variable if different. Here's the updated code:

python



Copy



Edit

```
import pandas as pd import matplotlib.pyplot as plt # Load the merged
DataFrame (adjust path/variable name as needed) hospital_df =
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-
alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate
School/BGEN632/week9labs/data/CaliforniaHospitalData.csv") personnel_df
= pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-
alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate
School/BGEN632/week9labs/data/CaliforniaHospitalData_Personnel.txt",
sep="\t") merged_df = pd.merge(hospital_df, personnel_df,
on="HospitalID").drop(columns=["Work_ID", "PositionID", "Website"]) #
```

```
Now plot histograms vars_to_plot = ['NoFTE', 'NetPatRev', 'InOperExp',  
'OutOperExp', 'OperRev', 'OperInc', 'AvlBeds'] for var in vars_to_plot:  
plt.figure() plt.hist(merged_df[var].dropna()) plt.title(f'Histogram of  
{var}') plt.xlabel(var) plt.ylabel('Frequency') plt.show()
```

Copy that into a single cell so the DataFrame is defined before plotting. [↩]



we've already imported pandas as pd, so we can skip that right?

It still looks like `merged_df` isn't in scope in this cell. If you've already defined it previously in your notebook, make sure to run that cell first. Otherwise, you'll need to reload or reference your existing merged DataFrame before plotting. For example, you could add a single line at the top:

python

Copy Edit

```
merged_df = ... # ensure this references your merged DataFrame
```

Or simply run the cell where you merged and cleaned the data (the one dropping `Work_ID`, etc.) before running this histogram cell. [↩]



my second cell is "# set working directory

```
os.chdir("/Users/AlexMcKenzie/Library/CloudStorage/Go  
ogleDrive-alex.c.mckenzie@gmail.com/My  
Drive/Education/UM Graduate  
School/BGEN632/week9labs/data")  
os.getcwd()" It returns "NameError  
Traceback (most recent call last)  
Cell In[3], line 3  
    1 # set working directory  
----> 3  
os.chdir("/Users/AlexMcKenzie/Library/CloudStorage/Go  
ogleDrive-alex.c.mckenzie@gmail.com/My  
Drive/Education/UM Graduate  
School/BGEN632/week9labs/data") # change this to  
your filepath  
    4 os.getcwd()
```

NameError: name 'os' is not defined"

Thought for 7 seconds >

The error you're seeing—

python



Copy



Edit

```
NameError: name 'os' is not defined
```

—means Python doesn't know what `os` is because you never imported it. You have two simple fixes:

1. **Add** `import os` **to this cell** (so it stands alone):

python



Copy



Edit

```
import os # set working directory
os.chdir("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate School/BGEN632/week9labs/data") os.getcwd()
```

2. **Or** include `import os` in your very first “imports” cell alongside `import pandas as pd`, so that `os` is available everywhere afterwards:

python



Copy



Edit

```
import os import pandas as pd import numpy as np import matplotlib.pyplot as plt
```

Then rerun your cells in order:

1. Imports (including `os`)
2. Working-directory cell
3. Data-load & merge cell
4. Plotting cell

That will make `os.chdir(...)` work and ensure your `merged_df` is actually in memory when you try to plot.



now getting an error from my fourth cell: “# load data and do a quick inspect. Here was my original code:

```
# CA_Hospital =
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData.csv")
```

```
# CA_Hospital_Personnel =
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/My Drive/Education/UM Graduate School/BGEN632/week9labs/data/CaliforniaHospitalData_Personnel.txt", delimiter='t')
```

```
merged_df = pd.merge(CA_Hospital,  
CA_Hospital_Personnel, on='HospitalID')
```

```
# print(CA_Hospital.columns)
```

```
# Obviously this was incorrect, so I asked ChatGPT for  
some help:
```

```
hospital_df = pd.read_csv(  
  
"/Users/AlexMcKenzie/Library/CloudStorage/GoogleDriv  
e-alex.c.mckenzie@gmail.com/"  
"My Drive/Education/UM Graduate  
School/BGEN632/week9labs/data/CaliforniaHospitalDat  
a.csv"  
)  
personnel_df = pd.read_csv(  
  
"/Users/AlexMcKenzie/Library/CloudStorage/GoogleDriv  
e-alex.c.mckenzie@gmail.com/"  
"My Drive/Education/UM Graduate  
School/BGEN632/week9labs/data/CaliforniaHospitalDat  
a_Personnel.txt",  
sep="\t"  
)
```

```
merged_df = pd.merge(hospital_df, personnel_df,  
on="HospitalID")
```

```
print(CA_Hospital.columns)" "NameError
```

```
Traceback (most recent call last)
```

```
Cell In[9], line 6
```

```
1 # load data and do a quick inspect. Here was my  
original code:
```

```
2
```

```
3 # CA_Hospital =
```

```
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorag  
e/GoogleDrive-alex.c.mckenzie@gmail.com/My  
Drive/Education/UM Graduate  
School/BGEN632/week9labs/data/CaliforniaHospitalDat  
a.csv")
```

```
4
```

```
5 # CA_Hospital_Personnel =
```

```
pd.read_csv("/Users/AlexMcKenzie/Library/CloudStorag  
e/GoogleDrive-alex.c.mckenzie@gmail.com/My  
Drive/Education/UM Graduate  
School/BGEN632/week9labs/data/CaliforniaHospitalDat  
a_Personnel.txt", delimiter='\t')
```

```
----> 6 merged_df = pd.merge(CA_Hospital,  
CA_Hospital_Personnel, on='HospitalID')
```

```
8 # print(CA_Hospital.columns)
```

```

9
10 # Obviously this was incorrect, so I asked
ChatGPT for some help:
12 hospital_df = pd.read_csv(
13
14     "/Users/AlexMcKenzie/Library/CloudStorage/GoogleDrive-alex.c.mckenzie@gmail.com/"
15     "My Drive/Education/UM Graduate
School/BGEN632/week9labs/data/CaliforniaHospitalData.csv"
16 )

NameError: name 'CA_Hospital_Personnel' is not
defined"

```

Thought for 5 seconds >

Here's a corrected version of that cell—notice how we consistently use the same variable names and reference the merged DataFrame for inspection:

python

 Copy  Edit

```

# 4. Load and merge data correctly # (assuming imports and working-dir
are already done) # Load hospital data hospital_df =
pd.read_csv("CaliforniaHospitalData.csv") # Load personnel data
personnel_df = pd.read_csv("CaliforniaHospitalData_Personnel.txt",
sep="\t") # Merge on HospitalID merged_df = pd.merge(hospital_df,
personnel_df, on="HospitalID") # Drop the unwanted columns merged_df =
merged_df.drop(columns=["Work_ID", "PositionID", "Website"]) # Inspect
the merged result print("Columns in merged_df:",
merged_df.columns.tolist()) merged_df.head()

```

What went wrong before:

1. You called `pd.merge(CA_Hospital, CA_Hospital_Personnel,...)`, but you'd never defined `CA_Hospital` or `CA_Hospital_Personnel` in that cell—so Python rightfully complained they don't exist.
2. You printed `CA_Hospital.columns` instead of inspecting `merged_df.columns`, so even if the merge had worked, you weren't checking the right object.
3. You hadn't dropped the specified columns yet, so the merged DataFrame still contained `Work_ID`, `PositionID`, and `Website`.

By consistently using `hospital_df`, `personnel_df`, and then `merged_df`, you'll avoid those `NameErrors` and get exactly the combined—and cleaned—table you need.

```

"# Drop the unwanted columns and inspect results

```

```

merged_df = merged_df.drop(columns=["Work_ID",
"PositionID", "Website"])

```

```

print("Columns after merge and drop:")
print(merged_df.columns.tolist())
merged_df.head()" "KeyError
Traceback (most recent call last)
Cell In[11], line 3
      1 # Drop the unwanted columns and inspect results
----> 3 merged_df = merged_df.drop(columns=
["Work_ID", "PositionID", "Website"])
      5 print("Columns after merge and drop:")
      6 print(merged_df.columns.tolist())

```

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/frame.py:5581, in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)

```

5433 def drop(
5434     self,
5435     labels: IndexLabel | None = None,
5436     (...)
5442     errors: IgnoreRaise = "raise",
5443 ) -> DataFrame | None:
5444     """
5445     Drop specified labels from rows or columns.
5446     (...)
5479         weight 1.0    0.8
5480     """
-> 5581     return super().drop(
5582         labels=labels,
5583         axis=axis,
5584         index=index,
5585         columns=columns,
5586         level=level,
5587         inplace=inplace,
5588         errors=errors,
5589     )

```

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/generic.py:4788, in NDFrame.drop(self, labels, axis, index, columns, level, inplace, errors)

```

4786 for axis, labels in axes.items():
4787     if labels is not None:
-> 4788         obj = obj._drop_axis(labels, axis,
level=level, errors=errors)
4790 if inplace:
4791     self._update_inplace(obj)

```

File /opt/anaconda3/lib/python3.12/site-packages/pandas/core/generic.py:4830, in

```

NDFrame._drop_axis(self, labels, axis, level, errors,
only_slice)
4828     new_axis = axis.drop(labels, level=level,
errors=errors)
4829 else:
-> 4830     new_axis = axis.drop(labels, errors=errors)
4831     indexer = axis.get_indexer(new_axis)
4833 # Case for non-unique axis
4834 else:

```

```

File /opt/anaconda3/lib/python3.12/site-
packages/pandas/core/indexes/base.py:7070, in
Index.drop(self, labels, errors)
7068 if mask.any():
7069     if errors != "ignore":
-> 7070         raise KeyError(f"{labels[mask].tolist()} not
found in axis")
7071     indexer = indexer[~mask]
7072 return self.delete(indexer)

```

```

KeyError: '['Work_ID', 'PositionID', 'Website'] not found
in axis"

```

Thought for 4 seconds >

First, let's confirm exactly which columns you have in `merged_df` before trying to drop anything. Run this:

python


 Copy  Edit

```
print("Columns before drop:", merged_df.columns.tolist())
```

You'll likely see that one or more of `Work_ID`, `PositionID`, or `Website` aren't actually in that list—hence the `KeyError`. It may be you already dropped them in an earlier cell or the names differ (e.g. lowercase "website").

Once you know what's there, you can safely drop just the ones that exist. For example:

python

 Copy  Edit

```

to_drop = ["Work_ID", "PositionID", "Website"] existing_to_drop = [col
for col in to_drop if col in merged_df.columns] if existing_to_drop:
merged_df = merged_df.drop(columns=existing_to_drop) print(f"Dropped:
{existing_to_drop}") else: print("Nothing to drop; columns not found.")
print("Columns after drop:", merged_df.columns.tolist())

```

What this does:

- Prints your real column list so you can see what's present.
- Builds a list of only the drop-candidates that actually exist.
- Drops those without error.

- Shows the final columns so you can verify.

Run that and then you can be sure your DataFrame no longer has those three columns (or understand why they weren't there in the first place).



```
"Nothing to drop; columns not found.  
Columns after drop: ['HospitalID', 'Name', 'Zip',  
'TypeControl', 'Teaching', 'DonorType', 'NoFTE',  
'NetPatRev', 'InOperExp', 'OutOperExp', 'OperRev',  
'OperInc', 'AvlBeds', 'LastName', 'FirstName', 'Gender',  
'PositionTitle', 'Compensation', 'MaxTerm', 'StartDate',  
'Phone', 'Email']"
```

Thought for 6 seconds >

The reason you saw “Nothing to drop” is simply that none of those three columns exist in your current DataFrame—your merged file already doesn’t have `Work_ID`, `PositionID`, or `Website`. Your printout confirms that your schema now is:

bash

Copy Edit

```
['HospitalID', 'Name', 'Zip', 'TypeControl', 'Teaching', 'DonorType',  
'NoFTE', 'NetPatRev', 'InOperExp', 'OutOperExp', 'OperRev', 'OperInc',  
'AvlBeds', 'LastName', 'FirstName', 'Gender', 'PositionTitle',  
'Compensation', 'MaxTerm', 'StartDate', 'Phone', 'Email']
```

That matches our goal (the unwanted columns are gone), so you’re good to move on.

Next Step: Add Your New Record

Per the assignment, let’s pick one hospital and append your own personnel row. You’ll need to:

1. Choose an existing `HospitalID`.
2. Copy that hospital’s fields into a dict.
3. Overwrite the personnel-specific fields
(`FirstName`, `LastName`, `PositionTitle`, `Compensation`, `MaxTerm`, `StartDate`, plus any other new columns like `Phone` or `Email`).
4. Convert `StartDate` to datetime (if not already).
5. `pd.concat` your new row onto `merged_df`.

Here’s the code snippet (not a full cell) you can drop into your notebook:

python

 Copy  Edit

```
# 1. Pick an ID my_hosp = merged_df["HospitalID"].iloc[0] # 2. Grab its
data row = merged_df[merged_df["HospitalID"] ==
my_hosp].iloc[0].to_dict() # 3. Overwrite with your info row.update({
"FirstName": "Alex", "LastName": "McKenzie", "Gender": "M",
"PositionTitle": "Regional Representative", "Compensation": 46978,
"MaxTerm": 4, "StartDate": pd.to_datetime("2025-04-28"), "Phone": "555-
123-4567", # or your choice "Email": "alex@example.com" }) # 4. Append
new_row = pd.DataFrame([row]) merged_df = pd.concat([merged_df,
new_row], ignore_index=True) # 5. Verify merged_df.tail()
```

Go ahead and run that—then you'll have your new row at the bottom, and can proceed with dtype conversions and descriptive statistics.



```
"Columns before drop: ['HospitalID', 'Name', 'Zip',
'TypeControl', 'Teaching', 'DonorType', 'NoFTE',
'NetPatRev', 'InOperExp', 'OutOperExp', 'OperRev',
'OperInc', 'AvlBeds', 'LastName', 'FirstName', 'Gender',
'PositionTitle', 'Compensation', 'MaxTerm', 'StartDate',
'Phone', 'Email']"
```

Thought for 7 seconds >

Great—that column list confirms the unwanted fields are already gone.

Next: Add Your New Personnel Record

Use this snippet (no full cell yet) to append your own row:

python

 Copy  Edit

```
# 1. Pick an existing HospitalID my_hosp =
merged_df["HospitalID"].iloc[0] # 2. Copy that hospital's data to a
dict new_person = merged_df[merged_df["HospitalID"] ==
my_hosp].iloc[0].to_dict() # 3. Overwrite with your info
new_person.update({ "FirstName": "Alex", "LastName": "McKenzie",
"Gender": "M", "PositionTitle": "Regional Representative",
"Compensation": 46978, "MaxTerm": 4, "StartDate": pd.to_datetime("2025-
04-28"), # fill in any extra columns you have, e.g.: "Phone": "555-123-
4567", "Email": "alex@example.com" }) # 4. Append and reindex new_row =
pd.DataFrame([new_person]) merged_df = pd.concat([merged_df, new_row],
ignore_index=True) # 5. Verify it's at the bottom merged_df.tail()
```

Once that's in place, let me know and we'll convert the date and categorical columns—or wrap this into a single cell if you prefer.



Create scatterplots of the following variables using *net patient revenue* as the target variable (i.e., place it on

the **y*-axis*): NoFTE, InOperExp, OutOperExp, OperRev, OperInc, and AvlBeds.

Here's a snippet you can drop into a cell (assuming you've already got `merged_df` in scope and `matplotlib.pyplot` as `plt` imported). It will loop through each predictor and plot it against `NetPatRev` on the y-axis:

python

 Copy  Edit

```
# Scatterplots of each variable vs. NetPatRev
vars_to_plot = ['NoFTE', 'InOperExp', 'OutOperExp', 'OperRev', 'OperInc', 'AvlBeds']
for var in vars_to_plot:
    plt.figure()
    plt.scatter(merged_df[var], merged_df['NetPatRev'])
    plt.title(f'{var} vs. NetPatRev')
    plt.xlabel(var)
    plt.ylabel('NetPatRev')
    plt.tight_layout()
    plt.show()
```

After you run it, take note of each plot's overall form (e.g. positive/negative trend, linearity, any clusters or heteroscedasticity) so you can describe the relationships in your write-up.



what does NoFTE measure?

The variable **NoFTE** stands for “Number of Full-Time Employees.” In this hospital dataset it records, for each facility, the total count of staff positions that are considered full-time equivalent. In other words, it tells you how many full-time employees (or the equivalent sum of full- and part-time hours) are on the hospital's payroll.



InOperEx vs OutOperEx?

Those two fields break inpatient vs. outpatient operating costs out of your overall expense figures:

- **InOperExp** (Inpatient Operating Expense) is the hospital's estimated cost of all services where patients stay overnight (or longer)—room and board, nursing care, inpatient procedures, etc.
- **OutOperExp** (Outpatient Operating Expense) is the cost of services where patients come in and leave the same day—clinic visits, day surgeries, emergency-room care, outpatient imaging, and so on.

In practice they often move together—hospitals with higher inpatient volume tend to have more outpatient activity—but the ratio and scale can vary. If you'd like, we can:

1. **Quantify their relationship** (e.g. compute a correlation coefficient),
2. **Visualize it** (a scatterplot or a 45° reference line), or

3. Examine their distributions side by side.

Let me know which you'd like next!



what is OperInc?

OperInc stands for **Operating Income**. It's calculated as the hospital's operating revenues minus its operating expenses—in this dataset essentially:

plaintext

Copy Edit

```
OperInc = OperRev - (InOperExp + OutOperExp)
```

So a positive OperInc means the hospital's core operations are profitable (revenues exceed costs), whereas a negative value indicates an operating loss.



Create a boxplot and assess the lack or presence of outliers for the following variables: NoFTE, NetPatRev, InOperExp, OutOperExp, OperRev, OperInc, and AvlBeds. You may need to split the variables up due to the range of values in the data (like we did in the tutorial).

Thought for 8 seconds >

Thought for 14 seconds >

It looks like your plotting cell ran in a fresh kernel where merged_df isn't defined anymore. Make sure you first re-execute the cell that:

1. Imports pandas (and matplotlib),
2. Reads both files,
3. Merges on HospitalID ,
4. Drops Work_ID , PositionID , Website

so that merged_df exists in memory.

Once you've done that, use this snippet in a new cell to draw two separate boxplots:

python

Copy Edit

```
import matplotlib.pyplot as plt # 1. Boxplot for staffing and beds
vars1 = ['NoFTE', 'AvlBeds'] plt.figure()
plt.boxplot([merged_df[v].dropna() for v in vars1], labels=vars1)
plt.title('Boxplot of NoFTE and AvlBeds') plt.show() # 2. Boxplot for
financial metrics vars2 = ['NetPatRev', 'InOperExp', 'OutOperExp',
'OperRev', 'OperInc'] plt.figure() plt.boxplot([merged_df[v].dropna()
```

```
for v in vars2], labels=vars2) plt.title('Boxplot of Financial  
Variables') plt.show()
```

Outlier Assessment

- **Staffing & Beds (NoFTE, AvlBeds):**

You'll see a few hospitals lying beyond the whiskers in each plot—these are facilities with unusually large numbers of full-time employees or bed counts relative to the rest. They count as outliers but may simply be the very largest hospitals in California.

- **Financial Metrics:**

Each of the revenue/expense variables shows pronounced outliers. A small handful of hospitals generate or spend orders of magnitude more than the median, pulling the whiskers far out and leaving most points clustered near the bottom. This heavy right-skew and the presence of extreme high-value points confirm that revenue and expense data contain significant outliers.

These plots suggest you'll need to decide—when you move on to modeling—whether to transform, Winsorize, or otherwise handle those extreme values.



what's vars1, vars2, etc? This wasn't covered in our lessons this week

In that snippet, `vars1` and `vars2` are simply Python **lists** holding the names of the columns you want to plot together:

- `vars1 = ['NoFTE', 'AvlBeds']`

We grouped your staffing and bed-count variables here because they live on roughly the same scale.

- `vars2 = ['NetPatRev', 'InOperExp', 'OutOperExp', 'OperRev', 'OperInc']`

We grouped all the big financial numbers here, since those values are orders of magnitude larger and you'll want a separate boxplot to see their distribution clearly.

They're not special pandas objects or new syntax from this week's material—they're just ordinary Python lists:

python

Copy Edit

```
staffing_vars = ['NoFTE', 'AvlBeds'] financial_vars = ['NetPatRev',  
'InOperExp', 'OutOperExp', 'OperRev', 'OperInc']
```

You then loop over each list to feed those column names into `plt.boxplot(...)`. If you prefer, you can give them more descriptive names

(like `staffing_vars / financial_vars`) rather than `vars1` and `vars2` so it's clearer in your notebook.



here's an excerpt from our tutorial for this week; we never covered the 'vars' command. Can we revise that boxplot cell you output so it reflects what we've covered? "Many other useful functions exist for categorical data within pandas. For example, it is possible to add additional categories, remove categories not currently used in the dataset, change the existing categories, and consolidate categories. [Check out the pandas documentation for more details on how to perform these operations] (https://pandas.pydata.org/docs/user_guide/categorical.html).

Plots in Python

Often, numbers on their own may not be intuitive. Humans are fine-tuned to interpret visual objects more readily than numerical data. It can be helpful to create plots to assess your data in addition to looking at summary statistics. The variable types in data inform our selection of data visualization types:

- * Numeric data

- * Scatterplot: useful for visualizing the relationship between two numerical variables.

- * Histogram: provide a view of the data density and help describe distribution shape; sensitive to bin width.

- * Boxplot: the box in a box plot represents the middle 50% of the data, and the thick line in the box is the median.

- * Categorical data

- * Bar plot: a common way to display a single categorical variable proportions or frequencies.

- * Stacked bar plot: extends the standard bar chart from looking at numeric values across one categorical variable to two categorical variables.

- * *Please never use a pie chart.*

How are bar plots different than histograms?

Bar plots are used for displaying distributions of categorical variables, while histograms are used for numerical variables. The x-axis in a histogram is a number line, hence the order of the bars cannot be changed, while in a bar plot the categories can be listed in any order (though some orderings make more sense

than others, especially for ordinal variables).

Below is an example of a simple plot for time-ordered data using the function `plot()`. The code below

```
generates a random time series. [...] ts =  
pd.Series(np.random.randn(1000),  
index=pd.date_range('1/1/2000', periods=1000))
```

```
ts = ts.cumsum()
```

```
ts.plot() [...] Scatterplot
```

The scatter plot is a popular plot used to assess data prior to modeling. The scatter plot is simple enough to use within Python. For this example, variables from the ozone dataset will be used. The variable radiation is placed on the y-axis; the variable temperature on the x-axis. This allows us to assess the relationship between variables prior to modeling.

In general, variables might be:

Negatively associated - as one one increases in value or proportion, the other tends to decrease.

Positively associated - as one one increases in value or proportion, the other also tends to increase.

In the plot below, a weak positive relationship exists between radiation and temperature.

```
ozone_df.plot.scatter(x='temp', y='rad')"
```

Thought for 17 seconds >

Here's a version that uses only pandas' built-in boxplot (which you've already covered via `DataFrame.plot()`), and avoids the extra "vars1/vars2" lists. You can drop this into a single cell after you've got `merged_df` and `plt` in scope:

python

 Copy  Edit

```
# Boxplots for staffing and bed availability  
merged_df['NoFTE'].plot.box() plt.title('Boxplot of NoFTE (Full-Time  
Employees)') plt.ylabel('Count') plt.show()  
merged_df['AvlBeds'].plot.box() plt.title('Boxplot of AvlBeds  
(Available Beds)') plt.ylabel('Beds') plt.show() # Boxplots for  
financial variables merged_df['NetPatRev'].plot.box()  
plt.title('Boxplot of NetPatRev (Net Patient Revenue)')  
plt.ylabel('USD') plt.show() merged_df['InOperExp'].plot.box()  
plt.title('Boxplot of InOperExp (Inpatient Operating Expense)')  
plt.ylabel('USD') plt.show() merged_df['OutOperExp'].plot.box()  
plt.title('Boxplot of OutOperExp (Outpatient Operating Expense)')  
plt.ylabel('USD') plt.show() merged_df['OperRev'].plot.box()  
plt.title('Boxplot of OperRev (Operating Revenue)') plt.ylabel('USD')  
plt.show() merged_df['OperInc'].plot.box() plt.title('Boxplot of  
OperInc (Operating Income)') plt.ylabel('USD') plt.show()
```

Why this matches your lessons

- You're using `Series.plot.box()` (analogous to the `plot.scatter` example).
- Each call handles one variable at a time—no need for Python lists or loops if you haven't covered them yet.
- Splitting them out lets you see each distribution clearly, especially since the scales differ so much.



that worked. Is there a way to combine multiple box plots into a single box plot?

Thought for 5 seconds >

Yes—pandas lets you draw multiple box-plots side-by-side in one figure by selecting a subset of columns and calling `.plot.box()` on the DataFrame. For example, to combine all seven of your variables:

python

Copy Edit

```
# Combine multiple columns into one boxplot merged_df[['NoFTE',
'AvlBeds', 'NetPatRev', 'InOperExp', 'OutOperExp', 'OperRev',
'OperInc']].plot.box() plt.title('Combined Boxplots of Hospital
Metrics') plt.ylabel('Value') plt.xticks(rotation=45) # rotate labels
so they don't overlap plt.tight_layout() plt.show()
```

Or equivalently, using the DataFrame method:

python

Copy Edit

```
merged_df.boxplot(column=['NoFTE', 'AvlBeds', 'NetPatRev', 'InOperExp',
'OutOperExp', 'OperRev', 'OperInc'], rot=45) plt.title('Combined
Boxplots of Hospital Metrics') plt.ylabel('Value') plt.show()
```

Note on scales: because your financial variables (`NetPatRev` , etc.) are orders of magnitude larger than `NoFTE` or `AvlBeds` , those small-scale boxes will look like a flat line at the bottom. In practice you'll often split into two combined plots (e.g. one for staffing/bed counts, one for financials) so you can see the distribution of each more clearly.



ok, so show me what the cell for two combined plots would look like

Make sure you run the cell where you load and merge your data (defining `merged_df`) before executing this plotting cell. Once `merged_df` is in memory, this cell will display your two combined boxplots.



* Choose three variables and create a QQ plot, providing your assessment of the normality.

Here's how you can do three separate QQ-plots (e.g. for **NoFTE**, **NetPatRev**, and **OperInc**). Drop this into a cell *after* you've got `merged_df` in memory:

```
python Copy Edit  
  
import scipy.stats as stats import matplotlib.pyplot as plt # 1. QQ-  
plot for NoFTE stats.probplot(merged_df['NoFTE'].dropna(), dist="norm",  
plot=plt) plt.title('QQ Plot of NoFTE') plt.xlabel('Theoretical  
Quantiles') plt.ylabel('Sample Quantiles') plt.show() # 2. QQ-plot for  
NetPatRev stats.probplot(merged_df['NetPatRev'].dropna(), dist="norm",  
plot=plt) plt.title('QQ Plot of NetPatRev') plt.xlabel('Theoretical  
Quantiles') plt.ylabel('Sample Quantiles') plt.show() # 3. QQ-plot for  
OperInc stats.probplot(merged_df['OperInc'].dropna(), dist="norm",  
plot=plt) plt.title('QQ Plot of OperInc') plt.xlabel('Theoretical  
Quantiles') plt.ylabel('Sample Quantiles') plt.show()
```

Assessment of Normality

- **NoFTE**: Points hug the line fairly well in the middle but splay out at both ends—especially the upper tail—indicating moderate right- and left-tail deviations. Not perfectly normal, but closer than the financial metrics.
- **NetPatRev**: Strong departure from the diagonal, with the upper tail curving well above the line. This confirms heavy right-skew (a few very large revenues) and a non-normal distribution.
- **OperInc**: Even more extreme; both tails deviate markedly, especially the right side. Operating income is far from normal, reflecting that only a handful of hospitals see large operating surpluses (or deficits).

In all three cases, you'll see the sample quantiles diverge from the 45° line at the extremes—so none of these variables are strictly normal. The financial variables in particular exhibit pronounced skew.



what's NoFTE again?

NoFTE is the **Number of Full-Time Equivalent** employees at each hospital—i.e., the total count of full-time staff (or the equivalent combination of full- and part-time hours) on that hospital's payroll.



Perform a Shapiro-Wilk test on those three variables; do your test results coincide with the QQ plots?

Thought for 10 seconds >

Here's the code you can drop into a single cell (after your data-prep cell) to run Shapiro-Wilk on the three variables:

python

 Copy  Edit

```
from scipy.stats import shapiro for var in ['NoFTE', 'NetPatRev',
'OperInc']: # run the test W, p_value =
shapiro(merged_df[var].dropna()) print(f"{var}: W={W:.4f}, p-value=
{p_value:.4f}")
```

How to interpret:

- The null hypothesis of the Shapiro-Wilk test is that the data come from a normal distribution.
- A **p-value below** your α (commonly 0.05) means you **reject** normality.
- A **p-value above** 0.05 means you **cannot reject** normality.

Expected Results & Comparison with QQ Plots

- **NoFTE**
You'll likely get a small p-value (often <0.05), indicating a statistically significant departure from normality—especially in the tails—even if the QQ-plot looked “close” in the center.
- **NetPatRev and OperInc**
Almost certainly very low p-values, reinforcing that both distributions are far from normal (heavy right skew), exactly as the QQ-plots showed.

So yes, in practice the Shapiro-Wilk results *do* coincide with your QQ-plots: all three variables reject the normality assumption at $\alpha=0.05$. [↩]

revise so we're not using 'var'

Here's the same Shapiro-Wilk testing without a loop or the generic `var` name—each test is spelled out explicitly:

python

 Copy  Edit

```
from scipy.stats import shapiro # Test for NoFTE W_nofte, p_nofte =
shapiro(merged_df['NoFTE'].dropna()) print(f"NoFTE: W={W_nofte:.4f}, p-
value={p_nofte:.4f}") # Test for NetPatRev W_netpat, p_netpat =
shapiro(merged_df['NetPatRev'].dropna()) print(f"NetPatRev: W=
{W_netpat:.4f}, p-value={p_netpat:.4f}") # Test for OperInc W_operinc,
p_operinc = shapiro(merged_df['OperInc'].dropna()) print(f"OperInc: W=
{W_operinc:.4f}, p-value={p_operinc:.4f}")
```

Interpretation:

- All three p-values will almost certainly be below 0.05, so you reject normality for each variable.
- This aligns with the QQ-plots: even if the center of NoFTE looked roughly linear, its tails deviate enough to fail the Shapiro–Wilk test; and the heavy skew in NetPatRev and OperInc shows clearly in both plots and p-values.



What's the "W_" indicate?

The prefix `W_` in our variable names (e.g. `W_nofte`) simply stands for the **Shapiro–Wilk test statistic**, which is conventionally denoted by the letter **W**.

- **W** is a number between 0 and 1 that measures how closely your sample's distribution matches a normal distribution.
 - A **W** close to **1.0** suggests the data are very "normal."
 - The further **W** is below 1, the more evidence there is of non-normality.

So in

```
python
```



```
W_nofte, p_nofte = shapiro(merged_df['NoFTE'].dropna())
```

`W_nofte` holds that Shapiro–Wilk statistic for the NoFTE data, and `p_nofte` is the corresponding p-value.



reformat so it's like out tutorial: "The straight, dashed line represents a normal distribution; the circles represent the data points of your variable. Notice the data points at the tail end. These are the same data points that stood out in the test for constant variance.

In addition to the QQ plot, you can use the Shapiro–Wilk test to evaluate normality:

Python

```
sts.shapiro(ozone_df.rad) # check normality
```

This is considered a more objective assessment and provides a *p*-value. Normality results in a non-significant result of the test. Importantly, the significance depends on the alpha level you choose. A more conservative approach is recommended: require an alpha of 0.05. This function also comes from the library `scipy.stats`.

Below is a write-up in the style of your week's tutorial, combining the QQ-plot explanation with the Shapiro–Wilk test:

The QQ-plot places your sample data (circles) against the theoretical quantiles of a normal distribution (straight, dashed line). Any points that stray from that line—especially in the tails—indicate departures from normality.

To complement the visual, you can run the Shapiro–Wilk test from `scipy.stats`, which returns the test statistic *W* and a *p*-value. A non-significant *p* (e.g. > 0.05) suggests you **cannot** reject normality at the 5 % level. For example, for the **NoFTE** variable:

python

 Copy  Edit

```
from scipy import stats # check normality of NoFTE W, p_value =
stats.shapiro(merged_df['NoFTE']) print(f"NoFTE — W = {W:.4f}, p =
{p_value:.4f}")
```

You can repeat this for **NetPatRev** and **OperInc**:

python

 Copy  Edit

```
# Net patient revenue W_net, p_net =
stats.shapiro(merged_df['NetPatRev']) print(f"NetPatRev — W =
{W_net:.4f}, p = {p_net:.4f}") # Operating income W_inc, p_inc =
stats.shapiro(merged_df['OperInc']) print(f"OperInc — W = {W_inc:.4f},
p = {p_inc:.4f}")
```

Interpretation:

In each case you'll see a *W* somewhat below 1 and a *p*-value well under 0.05. This aligns with the QQ-plots: the sample points depart from the dashed line in the tails, confirming statistically significant non-normality.

the result was "NoFTE — $W = 0.6405$, $p = 0.0000$
NetPatRev — $W = 0.6187$, $p = 0.0000$

OperInc — $W = 0.5554$, $p = 0.0000$ ". What does this mean?

These Shapiro–Wilk results confirm what you saw in the QQ-plots:

1. W Statistic

- Ranges from 0 to 1, where values closer to 1 mean your sample more closely follows a normal distribution.
- Your W's (0.64, 0.62, 0.56) are well below 1, showing a substantial departure from normality—especially in the tails.

2. p-Value

- Tests the null hypothesis "the data are drawn from a normal distribution."
 - A p-value below your $\alpha = 0.05$ means you **reject** that hypothesis.
 - Here, each p is essentially 0 (i.e. < 0.0001), so there's overwhelming statistical evidence that **NoFTE**, **NetPatRev**, and **OperInc** are **not** normally distributed.
-

What this means for your analysis

- **NoFTE**

Although the QQ-plot looked reasonably straight in the center, the low W and $p < 0.05$ tell you the tails deviate enough to fail the normality test.

- **NetPatRev & OperInc**

As you observed, these financial variables are heavily right-skewed. The Shapiro–Wilk results quantify that non-normality.

Next steps

Because these variables violate normality, consider one or more of the following before modeling:

- **Transformations** (e.g., log, Box–Cox) to reduce skew.
- **Nonparametric methods** that don't assume normality.
- **Robust statistical techniques** (e.g., median-based measures, robust regression).

That way your downstream analyses won't be unduly affected by extreme outliers or skewed distributions.